

NATIONAL UNIVERSITY OF SINGAPORE

CS2106 – INTRODUCTION TO OPERATING SYSTEM

(Semester 1: AY2015/16)

Time Allowed: 2 Hours

INSTRUCTIONS TO CANDIDATES

1. This assessment paper consists of **SIXTEEN (16)** questions and comprises **TEN (10)** printed pages.
2. This is a **CLOSED BOOK** assessment. Two handwritten A4 reference sheets are allowed. Calculators are not allowed.
3. Answer all questions and write your answers in the **ANSWER BOOKLET** provided.
4. Fill in your Student Number with a **pen, clearly on odd-numbered pages** of your **ANSWER BOOKLET**.
5. You may use pencil to write your answers.
6. Marks allocated to each question are indicated. Total marks for the paper is **100**.
7. You are to submit only the **ANSWER BOOKLET** and no other document.

Questions 1 - 10: Each question has only one correct answer. Write your answers in the boxes provided in the Answer Booklet. Two marks are awarded for a correct answer and no penalty for wrong answer.

1. Suppose a compiled executable is moved from a system with 4KB page size to a new compatible system but with 8KB page size. There is no other difference in the execution environment other than the page size. Which of the following statement(s) is/are TRUE?
 - i. The executable need to be recompiled due to the page size difference.
 - ii. The size of page table is now $\frac{1}{2}$ of the original page table size.
 - iii. The size of page table entry is now $\frac{1}{2}$ of the original page table entry size.
 - a. (i) only.
 - b. (ii) only.
 - c. (ii) and (iii) only.
 - d. (i) and (iii) only.
 - e. (i), (ii) and (iii).

2. Under the buddy system, which of the following statement regarding a block B with starting address **40** is/are TRUE?
 - i. The buddy block of B at size 4 has starting address 44.
 - ii. The buddy block of B at size 8 has starting address 32.
 - iii. There is no buddy block of B at size 16 because B is not valid at that size.
 - a. (i) only.
 - b. (ii) only.
 - c. (ii) and (iii) only.
 - d. (i) and (iii) only.
 - e. (i), (ii) and (iii).

3. Suppose the content of a file F is spread across N disk sectors on a hard disk. Each track on the hard disk contains S sectors in total. Which of the following gives the most accurate approximation of the minimum and maximum number of **disk seeks** required to access the **entire** file F?
 - a. Minimum = 1 seek, Maximum = S seeks.
 - b. Minimum = S seeks, Maximum = N seeks.
 - c. Minimum = N / S seeks, Maximum = S seeks.
 - d. Minimum = N / S seeks, Maximum = N seeks.
 - e. None of the above.

4. On a system with 1024 physical memory frames and a 32-bit virtual memory space, which of the following regarding the structure of page table is/are TRUE?
- Even when the physical frames are not fully occupied, the **inverted page table** contains 1024 entries.
 - If **direct paging** is used, the total number of page table entries for a process depends on the amount of virtual memory space used.
 - If **2-level paging** is used, the total number of page directory entries for a process depends on the amount of virtual memory space used.
- (i) only.
 - (ii) only.
 - (ii) and (iii) only.
 - (i) and (iii) only.
 - (i), (ii) and (iii).
5. OS can use the system wide "open file table" to:
- Allow two or more processes to share a file.
 - Ensure each processes have their own set of file descriptors.
 - Provide protection of file between processes.
- (i) only.
 - (ii) only.
 - (ii) and (iii) only.
 - (i) and (iii) only.
 - (i), (ii) and (iii).
6. Under the FAT file system, to locate a file with filepath `"/d1/d2/d3/f"`, how many disk accesses are required? Note the following assumptions:
- The first `"/"` represents the root directory.
 - All directory entries (including the root directory) are in disk (i.e. not cached) at the beginning. Directory entries for all directories fit in one disk block.
 - The file path is valid.
- 1
 - 4
 - 5
 - 6
 - None of the above

7. Which of the following statement regarding execution mode change (user mode \leftrightarrow kernel mode) is TRUE?
- If pure user thread is used, thread switching does not involve an execution mode change.
 - If pure kernel thread is used, thread switching does not involve an execution mode change.
 - Process termination does not involve an execution mode change.
 - Closing a file does not involve an execution mode change.
 - None of the above.
8. Which of the following statement(s) is/are TRUE regarding the **interval between time interrupt (ITI)** and **time quantum (TQ)**?
- Shorter ITI (i.e. more frequent timer interrupt) but with the same TQ can increase the responsiveness of processes.
 - Shorter ITI reduces the overall CPU time spent on processes.
 - Shorter TQ reduces the overall CPU time spent on processes.
- (i) only.
 - (ii) only.
 - (ii) and (iii) only.
 - (i) and (iii) only.
 - (i), (ii) and (iii) only.
9. Which of the following statement(s) regarding the **zombie process state in Unix systems** is/are TRUE ?
- All child process will transit to zombie process state upon normal termination.
 - A process will transit to zombie process state even if it was terminated due to exception (e.g. segmentation fault).
 - Zombie process state was created so that wait() or similar system call can be performed correctly.
- (i) only.
 - (ii) only.
 - (ii) and (iii) only.
 - (i) and (iii) only.
 - (i), (ii) and (iii).

10. Which of the following statement(s) regarding the "Limited Seat" solution for the dining philosopher problem (reproduced below for your reference) is/are TRUE?

```

seats = Semaphore( 4 );
fork = Sempahore(1)[5]; //array of 5 semaphores
                                //each with initial value = 1
void philosopher( int i ){
    while (TRUE){
        Think( );
        wait( seats );
        wait( fork[LEFT] );
        wait( fork[RIGHT] );
        Eat( );
        signal( fork[LEFT] );
        signal( fork[RIGHT] );
        signal( seats );
    }
}

```

- i. A philosopher task can still stuck on the "wait(fork[LEFT])" under this solution.
 - ii. There could be up to 4 philosophers performing the "Eat()" at the same time.
 - iii. A philosopher task can still stuck on the "wait(fork[RIGHT])" under this solution.
- a. (i) only.
 - b. (ii) only.
 - c. (ii) and (iii) only.
 - d. (i) and (iii) only.
 - e. (i), (ii) and (iii) only.

11. [10 marks] Suppose we have two types of tasks on a system: **Coffee** and **Milk**. There are many tasks of each type running at any point in time. If we have a special “critical section” with the following requirement:

- Only a **pair of 1 Coffee and 1 Milk** task can enter the critical section at any time.
- Both Coffee and Milk must enter the critical section at the same time. (i.e. a single coffee task alone or a single milk task alone cannot be allowed to enter).

One possible solution is to use an additional task called the **Mixer** for monitoring the entry condition.

- a. [2 marks] Give the initialization value for each of the semaphore. You are allow to use **at most 4 semaphores** to solve this problem.
- b. [8marks] Give the implementation of the following 3 functions:
 - i. *CoffeeEnterCS()*: Coffee tasks call this when reaching the critical section.
 - ii. *MilkEnterCS()*: Milk tasks call this when reaching the critical section.
 - iii. *MixerMonitor()*: Mixer task run this “forever” to provide correct access to the critical section.

Important notes

- All functions are just a series of **wait()** / **signal()** calls. **You are not allowed to use any other programming construct in the functions.**
- Your task is just to ensure the entry condition of the "special critical section" is met. There is no need to consider how/when the coffee/milk task exit the critical section.

12. [16 marks]

- a. [4 marks] During the execution of a process, the OS may still take over the CPU sometimes. Give one example for each of the following scenarios.

- i. The running process explicitly invokes the OS to take over.
- ii. The OS takes over independent of the current running process.

- b. [6 marks] Consider the following pure CPU tasks and arrival time:

Task	Arrival Time	Total CPU time needed
A	0	3
B	0	4
C	3	4
D	5	2

Give the CPU schedule under Round Robin Algorithm with a time quantum of 2 time unit. Just indicate the tasks in the order of their execution.

You can assume: task arrives in the ready queue before the scheduler is triggered for that time unit. (e.g. at time 0, scheduler will have task A and B in ready queue for consideration).

- c. [6 marks] Given the following definition (all units are in *seconds*):

Definitions
Tcs: Time to context switch between two processes.
ITI: Interval between timer interrupt.
TQ: Time quantum.
Tsch: Time to perform scheduling bookkeeping, excluding the cost for actual context switch.

Express the overhead of scheduling per second of the **round robin scheduling algorithm**.

13. [10 marks]

- a. [6 marks] Given 3 memory frames, consider the following memory reference strings:

A	1	2	3	4	4	3	2	5	6	2	3
B	6	3	4	2	3	4	2	1	5	3	4

Give the number of page faults for OPT, LRU and Second Chance page replacement algorithms.

- b. [4 marks] Give 1 advantage and 1 disadvantage of the Second Chance algorithm as compared to LRU algorithm.

14. [16 marks] We can protect a memory page by adding permission bits to the page table entry (PTE). Suppose we add 3 bits: {**R**: Readable, **W**: Writable, **X**: Executable} to each PTE, answer the following briefly:

- [2 marks] What type of process instruction requires the checking of the "W" bit?
- [2 marks] If processor issues a "fetch instruction from page X", what are the permission bits to be checked?

When a processor instruction violates the access permission of a page, OS will be invoked to handle the problem. We can utilize this behaviour to implement the **copy-on-write** mechanism (memory pages are shared between parent and child until written).

Suppose process P has only 3 valid page table entries:

Page No	Frame No	R	W	X
0	7	1	0	1
1	2	1	1	0
2	5	1	0	0
3 N-1	---	---	---	---

- c. [4 marks] Give the page table entries for the **child process** after P executed a **fork()** system call. If you need to use any new frame numbers, use them in this order {6, 0, 3, 4 1}.

Using the above scenario, explain the following clearly:

- [2 marks] How does the OS know copy-on-write is needed?
- [6 marks] What are the steps required to handle copy-on-write? Indicate any additional information that OS need to maintain. Show the affected PTE(s) for the child process afterwards.

15. [12 marks]

- a. [4 marks] Give the disk I/O schedule for the following I/O requests using **Shortest Seek First (SSF)** algorithm:


Time	0	1	1	1	2	2	3	4
Request	19	2	6	9	11	17	4	5

For simplicity, we assume that all disk I/O requests can be completed in one time unit. An I/O request will be considered by the algorithm as soon as it arrives.

- b. [4 marks] Rearrange the above requests (both timing and order) such that it gives the **same schedule** under the C-SCAN algorithm. We assume the seek direction goes from track 1 to track 20. The requests should be issued as early as possible.
- c. [4 marks] Similarly, rearrange the above requests such that it gives the same schedule under the SCAN algorithm. The seek direction goes from track 1 to track 20, then track 20 back to track 1. The requests should be issued as early as possible.

16. [16 marks] Use the partial ext2 file system snapshot on the last page to answer the following questions. Note that only the crucial information are shown. **The "/" directory has the I-node number 8.**

- a. [4 marks] Give the sequence of **I-Node** accessed to "list the content" of directory **"/Zeta"**.
- b. [4 marks] Give the sequence of **data blocks** which stores the content for the file **"/Zeta/Alpha"**.
- c. [8 marks] Follow either of the following examples to draw / describe the entire directory structure **starting from the "/" directory**. Note that not all information given in the snapshot is linked to "/" directory.
- d.



I-Node Table		Disk Blocks	
1	9 7	1	5 F 5 Gamma 2 F 4 Beta
2	5 10	2	4 F 5 Alpha 3 D 5 Gamma
3	12	3	3 D 5 Gamma 1 F 5 Alpha
4	1 2	4	2 F 6 Lambda 4 F 5 Delta
5	3	5	7 F 4 Zeta 3 F 5 Gamma
6	4	6	5 D 4 Zeta 6 D 4 Beta
7	13 14	7	2 F 5 Sigma 7 F 4 Beta
8	6	8	1 F 5 Sigma 6 D 5 Gamma 4 D 5 Lambda
9	9	5 F 5 Gamma 2 F 4 Beta
		10	2 F 5 Sigma
		11	3 D 5 Gamma 1 F 5 Alpha
		12	2 F 5 Sigma
		13	4 F 5 Alpha
		14	5 D 4 Zeta 6 D 4 Beta
	

~~~~~ END OF PAPER ~~~~~